

Introduction to Deep Learning

Autoencoders

Andres Mendez-Vazquez

June 22, 2025

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part

3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

Outline

1 The Idea of The Autoencoder

● Introduction

- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part


3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

Introduction

From [1]

- “An autoencoder is a specific type of a neural network, which is mainly designed to encode the input into a compressed and meaningful representation, and then decode it back such that the reconstructed input is similar as possible to the original one.”



At the following work [2]

Definition

- An $n/p/n$ autonecoder is defined as a t -tuple $n, p, m, \mathbb{F}, \mathbb{G}, \mathcal{A}, \mathcal{B}, \mathcal{X}, \Delta$ where
 - 1 \mathbb{F} and \mathbb{G} are sets.
 - 2 n and p are positive integers with $0 < p < n$.
 - 3 \mathcal{A} is a class of functions from \mathbb{G}^p to \mathbb{F}^n .
 - 4 \mathcal{B} is a class of functions from \mathbb{F}^n to \mathbb{G}^p .
 - 5 $\mathcal{X} = \{x_1, \dots, x_m\}$ is a set of m (training) vectors in \mathbb{F}^n .
 - 6 Δ is a dissimilarity or distortion function over \mathbb{F}^n .

Basically

We have that

- For any $A \in \mathcal{A}$ and $B \in \mathcal{B}$, the autoencoder transforms an input vector $x \in F^n$ into an output vector $A \circ B(x) \in F^n$

Basically

We have that

- For any $A \in \mathcal{A}$ and $B \in \mathcal{B}$, the autoencoder transforms an input vector $x \in F^n$ into an output vector $A \circ B(x) \in F^n$

Thus, we have

- The corresponding autoencoder problem is to find $A \in \mathcal{A}$ and $B \in \mathcal{B}$, that minimize the overall distortion function:

$$\min E(A, B) = \min_{A, B} \sum_{t=1}^m E(x_t) = \min_{A, B} \sum_{t=1}^m \Delta(B \circ A(x_t), x_t)$$

And there is the other case

We can have another output y_t

$$\min_{A,B} \sum_{t=1}^m \Delta(B \circ A(x_t), y_t)$$

And there is the other case

We can have another output y_t

$$\min_{A,B} \sum_{t=1}^m \Delta(B \circ A(x_t), y_t)$$

Not only that

- $p < n$ corresponds to the regime where the autoencoder tries to implement some form of compression or feature extraction.

In 1989, they proposed an initial linear case no activation functions

$$E(A, B) = \sum_{1 \leq t \leq T} \|y_t - BAx_t\|$$

In 1989, they proposed an initial linear case no activation functions

$$E(A, B) = \sum_{1 \leq t \leq T} \|y_t - BAx_t\|$$

There are several properties under these autoencoders

- But somethin interesting an optimal:
 - ▶ $E(A, B)$ is convex in the coefficient of B and attains minimum for B such that $BAA^T \Sigma_{XX} = A^T \Sigma_{XY}$

Outline

1 The Idea of The Autoencoder

- Introduction
- **The Bottleneck Idea**
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part


3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

Basically, we have

Encoder


- The encoder is the part of the network which takes in the input and produces a lower Dimensional encoding



Basically, we have

Encoder

- The encoder is the part of the network which takes in the input and produces a lower Dimensional encoding



Something Notable


Bottleneck: It is the lower dimensional hidden layer where the encoding is produced.

- Note: Call it an embedding!!!

Then, we have that

Decoder


- The decoder takes in the encoding and recreates back the input.



Then, we have that

Decoder

- The decoder takes in the encoding and recreates back the input.

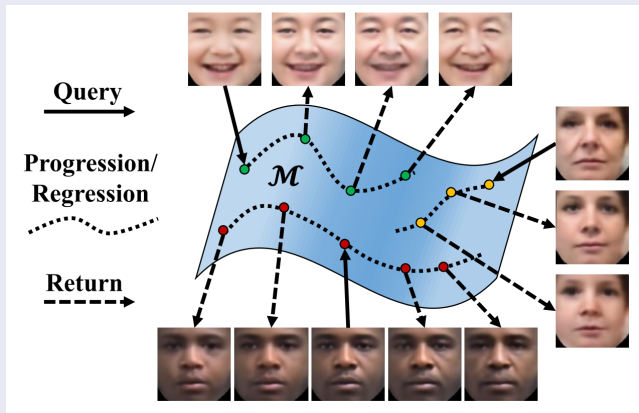


Remember the Transpose Convolution

- Yes, we can use it to implement the decoder.

Basically is Called a Manifold

Mapping x to a small dimension h from old to young



- Thanks to the Pytorch implementation by Mattan Serry, Hila Balahsan, and Dor Alt.

Autoencoders differ from General Data Compression

Autoencoders are data-specific

- i.e., only able to compress data similar to what they have been trained on

Autoencoders differ from General Data Compression

Autoencoders are data-specific

- i.e., only able to compress data similar to what they have been trained on

This is different from, say, MP3 or JPEG compression algorithm

- Which make general assumptions about "sound/images", but not about specific types of sounds/images
- Autoencoder for pictures of cats would do poorly in compressing pictures of trees
 - ▶ Because features it would learn would be cat-specific

Autoencoders differ from General Data Compression

Autoencoders are data-specific

- i.e., only able to compress data similar to what they have been trained on

This is different from, say, MP3 or JPEG compression algorithm

- Which make general assumptions about "sound/images", but not about specific types of sounds/images
- Autoencoder for pictures of cats would do poorly in compressing pictures of trees
 - ▶ Because features it would learn would be cat-specific

Autoencoders are lossy

- which means that the decompressed outputs will be degraded compared to the original inputs (similar to MP3 or JPEG compression).
- This differs from lossless arithmetic compression

And Actually, Autoencoders are Learn

Learning $g(f(x)) = x$ everywhere is not useful

- Actually in Age Autoencoders we want old faces from young

And Actually, Autoencoders are Learn

Learning $g(f(x)) = x$ everywhere is not useful

- Actually in Age Autoencoders we want old faces from young

Autoencoders are designed to be unable to copy perfectly

- Autoencoders learn useful properties of the data

And Actually, Autoencoders are Learn

Learning $g(f(x)) = x$ everywhere is not useful

- Actually in Age Autoencoders we want old faces from young

Autoencoders are designed to be unable to copy perfectly

- Autoencoders learn useful properties of the data

Autoencoders learn useful properties of the data

- Being forced to prioritize which aspects of input should be copied

An important property

It can learn stochastic mappings

- Go beyond deterministic functions to mappings $p_{\text{encoder}}(h|x)$ and $p_{\text{decoder}}(x|h)$

An important property

It can learn stochastic mappings

- Go beyond deterministic functions to mappings $p_{\text{encoder}}(h|x)$ and $p_{\text{decoder}}(x|h)$

Now

- A little bit on the loss function...

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- **Training Autoencoders**
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part

3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

Autoencoder is a feed-forward non-recurrent neural net

With an input layer, an output layer and one or more hidden layers

- Can be trained using the same techniques
- Compute gradients using back-propagation
 - ▶ Followed by minibatch gradient descent

Autoencoder is a feed-forward non-recurrent neural net

With an input layer, an output layer and one or more hidden layers

- Can be trained using the same techniques
- Compute gradients using back-propagation
 - ▶ Followed by minibatch gradient descent

Unlike feedforward networks, can also be trained using Recirculation

- Compare activations on the input to activations of the reconstructed input
- More biologically plausible than back-prop but rarely used in ML

Autoencoder training using a loss function

We have Encoder f and Decoder g

- $f : X \rightarrow h$
- $g : h \rightarrow X$
- Thus, we have

$$\arg \min_{f,g} \|x - (f \circ g) x\|^2$$

Example

One hidden layer

- Takes input $\mathbf{x} \in \mathbb{R}^d$
- Maps into an output $\mathbf{h} \in \mathbb{R}^p$

Example

One hidden layer

- Takes input $\mathbf{x} \in \mathbb{R}^d$
- Maps into an output $\mathbf{h} \in \mathbb{R}^p$

We have using an element wise activation function

- $\mathbf{h} = \sigma_1 (W\mathbf{x} + \mathbf{b}) \rightarrow \mathbf{x}' = \sigma_2 (W'\mathbf{h} + \mathbf{b}')$

Example

One hidden layer

- Takes input $\mathbf{x} \in \mathbb{R}^d$
- Maps into an output $\mathbf{h} \in \mathbb{R}^p$

We have using an element wise activation function

- $\mathbf{h} = \sigma_1 (W\mathbf{x} + \mathbf{b}) \rightarrow \mathbf{x}' = \sigma_2 (W'\mathbf{h} + \mathbf{b}')$

Trained to minimize reconstruction error

$$L(\mathbf{x}, \mathbf{x}') = \|\mathbf{x} - \mathbf{x}'\| = \left\| \mathbf{x} - \sigma_2 \left(W^t \sigma_1 (W\mathbf{x} + \mathbf{b}) + \mathbf{b}' \right) \right\|$$

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- **Encoder/Decoder Capacity**
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part

3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

Undercomplete Autoencoder

Copying input to output sounds useless

- Imagine that we do not have interest in decoder output
- We hope h takes on useful properties

Undercomplete Autoencoder

Copying input to output sounds useless

- Imagine that we do not have interest in decoder output
- We hope h takes on useful properties

Undercomplete autoencoder

- Constrain h to have lower dimension than x
- Force it to capture most salient features of training data

Autoencoder with linear decoder + MSE is PCA

We have that

$$L(\mathbf{x}, g(f(\mathbf{x}))) = \|\mathbf{x} - g(f(\mathbf{x}))\|^2$$

Autoencoder with linear decoder + MSE is PCA

We have that

$$L(x, g(f(x))) = \|x - g(f(x))\|^2$$

Something Notable

- Where L is a loss function penalizing $g(f(x))$ for being dissimilar from x

Autoencoder with linear decoder + MSE is PCA

We have that

$$L(\mathbf{x}, g(f(\mathbf{x}))) = \|\mathbf{x} - g(f(\mathbf{x}))\|^2$$

Something Notable

- Where L is a loss function penalizing $g(f(\mathbf{x}))$ for being dissimilar from \mathbf{x}

Thus

- When the decoder g is linear and L is the mean squared error, an undercomplete autoencoder learns to span the same subspace as PCA

$$g(\mathbf{x}) = W\mathbf{h} + \mathbf{b}$$

We have

In this case

- The autoencoder trained to perform the copying task has learned the principal subspace of the training data as a side-effect

We have

In this case

- The autoencoder trained to perform the copying task has learned the principal subspace of the training data as a side-effect

Autoencoders with nonlinear f and g

- They can learn more powerful nonlinear generalizations of PCA
- Something called Capacity

The capacity of a network

The capacity of a network refers to

- The range or scope of the types of functions that the model can approximate.
- Informally, a model's capacity is its ability to fit a wide variety of functions.

Encoder/Decoder Capacity

If encoder f and decoder g are allowed too much capacity

- Autoencoder can learn to perform the copying task without learning any useful information about distribution of data

Encoder/Decoder Capacity

If encoder f and decoder g are allowed too much capacity

- Autoencoder can learn to perform the copying task without learning any useful information about distribution of data

Something Notable

- Autoencoder with a one-dimensional code and a very powerful nonlinear encoder can learn to map $x(i)$ to code i .
- The decoder can learn to map these integer indices back to the values of specific training examples

Encoder/Decoder Capacity

If encoder f and decoder g are allowed too much capacity

- Autoencoder can learn to perform the copying task without learning any useful information about distribution of data

Something Notable

- Autoencoder with a one-dimensional code and a very powerful nonlinear encoder can learn to map $x(i)$ to code i .
- The decoder can learn to map these integer indices back to the values of specific training examples

A notable problem

- Autoencoder trained for copying task fails to learn anything useful if f/g capacity is too great

Cases when Autoencoder Learning Fails

Where autoencoders fail to learn anything useful

- Capacity of encoder/decoder f/g is too high
 - ▶ Capacity controlled by depth
- Hidden code \mathbf{h} has dimension equal to input \mathbf{x}
- Overcomplete case: where hidden code \mathbf{h} has dimension greater than input \mathbf{x}
 - ▶ Even a linear encoder/decoder can learn to copy input to output without learning anything useful about data distribution

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- **Right Autoencoder Design: Use regularization**
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part

3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

What is regularization

L1 Regularization

- This method adds a penalty to the loss function for the sum of the absolute values of the model weights.

What is regularization

L1 Regularization

- This method adds a penalty to the loss function for the sum of the absolute values of the model weights.

L2 Regularization

- This method adds a penalty to the loss function for the sum of the squares of the model weights.

What is regularization

L1 Regularization

- This method adds a penalty to the loss function for the sum of the absolute values of the model weights.

L2 Regularization

- This method adds a penalty to the loss function for the sum of the squares of the model weights.

Dropout

- This method randomly sets a fraction of the model's activations to zero during each training iteration.

Use regularization

Ideally

- choose code size (dimension of \mathbf{h}) small and capacity of encoder f and decoder g based on complexity of distribution modeled

Use regularization

Ideally

- choose code size (dimension of \mathbf{h}) small and capacity of encoder f and decoder g based on complexity of distribution modeled

Regularized autoencoders provide the ability to do so

- Rather than limiting model capacity by keeping encoder/decoder shallow and code size small
- They use a loss function that encourages the model to have properties other than copy its input to output=

Regularized Autoencoder Properties

Regularized AEs have properties beyond copying input to output

- Sparsity of representation
- Smallness of the derivative of the representation
- Robustness to noise
- Robustness to missing inputs

Regularized Autoencoder Properties

Regularized AEs have properties beyond copying input to output

- Sparsity of representation
- Smallness of the derivative of the representation
- Robustness to noise
- Robustness to missing inputs

Regularized autoencoder can be nonlinear and overcomplete

- But still learn something useful about the data distribution even if model capacity is great enough to learn trivial identity function

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- **Autoencoders as an initialization method**

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part

3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

Initializing Deep Learners

Autoencoders have many interesting applications

- As data compression, visualization, etc

Initializing Deep Learners

Autoencoders have many interesting applications

- As data compression, visualization, etc

Something Notable

- Y. Bengio, P. Lamblin, D. Popovici, and H. Larochelle. Greedy layer-wise training of deep networks. In Advances in Neural Information Processing Systems, 2007.

They discovered something interesting

Something Notable

- They observed that autoencoders could be used as a way to “pre-train” neural networks.

They discovered something interesting

Something Notable

- They observed that autoencoders could be used as a way to “pre-train” neural networks.

With pretraining, the process of training a deep network is divided in a sequence of steps

- 1 Pretraining step: train a sequence of shallow autoencoders, greedily one layer at a time, using unsupervised data,
- 2 Fine-tuning step 1: train the last layer using supervised data,
- 3 Fine-tuning step 2: use backpropagation to fine-tune the entire network using supervised data.

Basically

The step 2 can be seen as the training of a Perceptron

- This can be done easily

Basically

The step 2 can be seen as the training of a Perceptron


- This can be done easily

The Backpropagation of the entire network

- Also the classic procedure


The First Step is the interesting one

We have a Feed Forward Architecture




First Pre-training, the first layer

We have




Therefore, we have the final sub-step

We use the y to train another auto-encoder



Finally

We have that



However

As Quality Data Sets become more widely available

- Google, Facebook, Microsoft and many others took charge of that...

However

As Quality Data Sets become more widely available

- Google, Facebook, Microsoft and many others took charge of that...

This efforts were abandoned

- Instead Transfer Learning [4] become a more popular idea
 - ▶ For initialization....

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- **Sparse Autoencoders**
- Denoising Autoencoders
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part

3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

Sparse Autoencoders

Similar to the LASSO

$$\arg \min_{A,B} \sum_{t=1}^m \Delta(B \circ A(x_t), x_t) + \lambda \sum_i |a_i|$$

- where a_i is the activation at the i^{th} hidden layer and i iterates over all the hidden activation's.

Other Ways

KL-divergence

- We can assume the activation of each neuron acts as a Bernouli variable with probability p and tweak that probability.

Other Ways

KL-divergence

- We can assume the activation of each neuron acts as a Bernouli variable with probability p and tweak that probability.

For each neuron j

- The calculated empirical probability is $\hat{p}_j = \frac{1}{m} \sum_i a_i(i)$ where i iterates over the samples in the batch:

$$\arg \min_{A,B} \sum_{t=1}^m \Delta(B \circ A(x_t), x_t) + \sum_i KL(p \parallel \hat{p})$$

- ▶ where the regularization term in it aims at matching p to \hat{p} .

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- **Denoising Autoencoders**
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part


3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

Denoising Autoencoders

Note

- They can be viewed either as a regularization option, or as robust autoencoders which can be used for error correction.



Here, we have

We have that the input can be seen as

$$\tilde{x} = x + N(0, \sigma I) \rightarrow p(\tilde{x}|x) \sim N(x, \sigma I)$$

Here, we have

We have that the input can be seen as

$$\tilde{x} = x + N(0, \sigma I) \rightarrow p(\tilde{x}|x) \sim N(x, \sigma I)$$

An also

$$\tilde{x} = \beta \odot x \rightarrow \beta \sim Ber(p)$$

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- **Contractive Autoencoders**
- Example, Architecture of the U-Net
- Encoder Part

3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

Contractive Autoencoders

Here, the desire is to reduce the effect of small perturbations

- On the feature extraction process.

Contractive Autoencoders

Here, the desire is to reduce the effect of small perturbations

- On the feature extraction process.

By forcing the encoder to avoid changes

- Not important for the reconstruction by the decoder.

Basically

Apply regularization into the hidden layer of the encoder

- Into the Jacobian matrix of the Hidden Layers of the Encoder

Basically

Apply regularization into the hidden layer of the encoder

- Into the Jacobian matrix of the Hidden Layers of the Encoder

Formally, $J_{ji} = \nabla_{x_i} h_j(x_i)$

- Thus changes at the nodes h_j of a layer h .

Thus, we have

We try to minimize a Ridge regression on the Jacobian

$$\arg \min_{A,B} E(\Delta(x, B \circ A(x))) + \lambda \|J_A(x)\|_2^2$$

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- **Example, Architecture of the U-Net**
 - Encoder Part

3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

The Paper

Around 2015 [5]

- An extraordinary network came to be
 - ▶ “U-Net: Convolutional Networks for Biomedical Image Segmentation”
by Olaf Ronneberger, Philipp Fischer, and Thomas Brox

The Paper

Around 2015 [5]

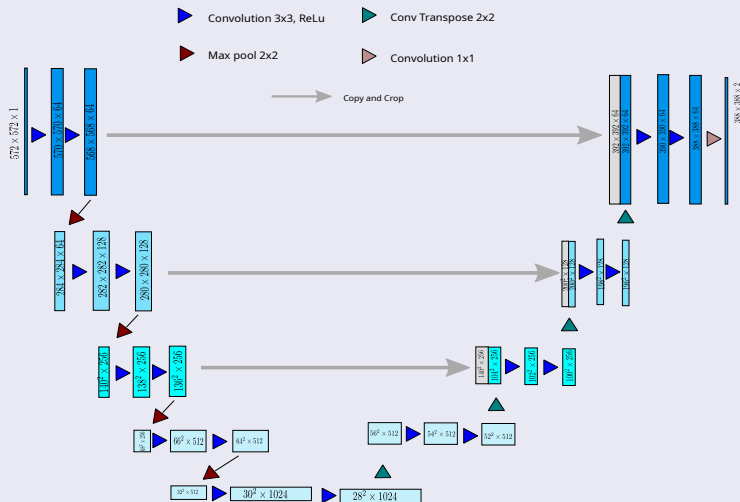
- An extraordinary network came to be
 - ▶ “U-Net: Convolutional Networks for Biomedical Image Segmentation”
by Olaf Ronneberger, Philipp Fischer, and Thomas Brox

It won

- 1 The Grand Challenge for Computer-Automated Detection of Caries in Bitewing Radiography at ISBI 2015,
- 2 The Cell Tracking Challenge at ISBI 2015

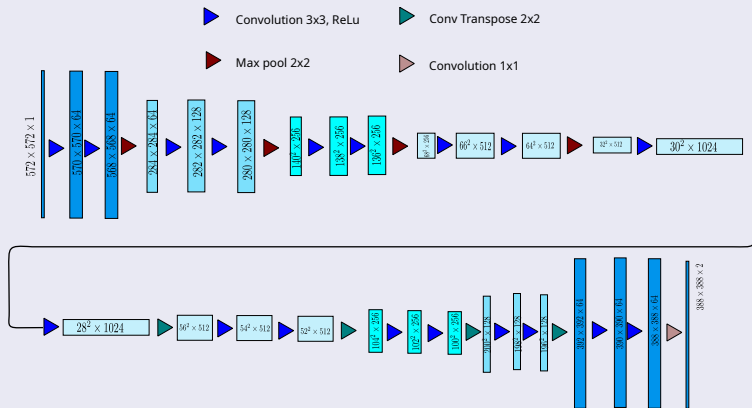
The Segmentation Architecture

We have (Imagen from [5])



Similarity with an Autoencoder

If we remove the extra passing of information, we finish with an autoencoder



Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

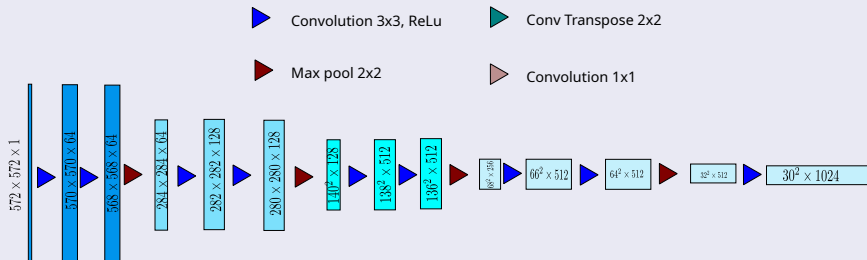
- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- **Example, Architecture of the U-Net**
 - **Encoder Part**

3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

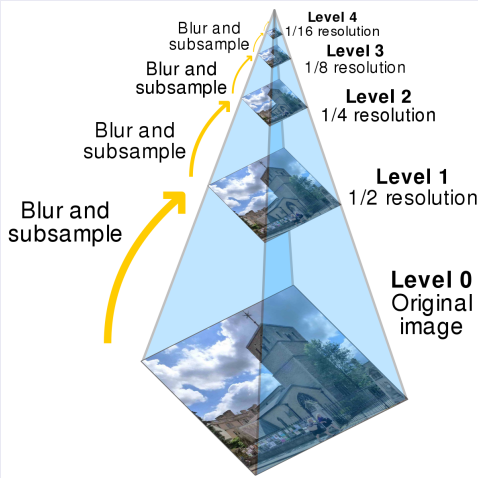
We have

At the encoder part



If you think about this

It is our Image Pyramid of old computer vision



Clearly

We can train this pyramide/encoder with data

- Yes our old backpropagation a.k.a automatic differentiation

Clearly

We can train this pyramid/encoder with data

- Yes our old backpropagation a.k.a automatic differentiation

At the center of this is the Convolution for multiple filters

$$Y_i^{(l)}(x, y) = B_i^{(l)}(x, y) + \sum_{j=1}^{m_1^{(l-1)}} \sum_{u=-ks}^{ks} \sum_{v=-ks}^{ks} Y_j^{(l-1)}(i-u, j-v) K_{ij}^{(l)}(u, v)$$

Basically

We take the j filter of the C_{out} dimension

- Then for each filter j you apply the convolution to each $Image_k$

Basically

We take the j filter of the C_{out} dimension

- Then for each filter j you apply the convolution to each $Image_k$

Then you simply add the images result of such convolution

- And Add Pointwise the bias $b_{C_{out_j}}$

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part

3 Variational Autoencoders

- **Introduction**
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

The Variational Autoencoders [6]

It is based on the Maximum a Posteriori Idea under a Random Process

- A value $x^{(i)}$ is generated from some likelihood distribution $p_{\Theta}(x|z)$ with a prior $p_{\Theta}(z)$

$$p_{\Theta}(z|x) \approx p_{\Theta}(x|z) p_{\Theta}(z)$$

The Variational Autoencoders [6]

It is based on the Maximum a Posteriori Idea under a Random Process

- A value $x^{(i)}$ is generated from some likelihood distribution $p_{\Theta}(x|z)$ with a prior $p_{\Theta}(z)$

$$p_{\Theta}(z|x) \approx p_{\Theta}(x|z) p_{\Theta}(z)$$

We ask $p_{\Theta}(x|z)$ and $p_{\Theta}(z)$

- 1 They are coming from parametric families
- 2 They are differentiable almost everywhere w.r.t. Θ and z

However

When we have the following case

- The case where the integral of the marginal likelihood is intractable

$$p_{\Theta}(x) = \int p_{\Theta}(x|z) p_{\Theta}(z) dz$$

However

When we have the following case

- The case where the integral of the marginal likelihood is intractable

$$p_{\Theta}(x) = \int p_{\Theta}(x|z) p_{\Theta}(z) dz$$

Also the **posterior density** is also intractable (No Expectation Maximization can be used)

$$p_{\Theta}(z|x) = \frac{p_{\Theta}(x|z) p_{\Theta}(z)}{p_{\Theta}(x)}$$

However

When we have the following case

- The case where the integral of the marginal likelihood is intractable

$$p_{\Theta}(x) = \int p_{\Theta}(x|z) p_{\Theta}(z) dz$$

Also the **posterior density** is also intractable (No Expectation Maximization can be used)

$$p_{\Theta}(z|x) = \frac{p_{\Theta}(x|z) p_{\Theta}(z)}{p_{\Theta}(x)}$$

These intractability's are quite common and appear in cases of moderately complicated likelihood functions

- For example Neural Networks...

More problems with sampling Bayesian Estimators

Gibbs Samplers, Metropolis-Hastings

- Have a serious problem, they are difficult to parallelize making them unusable for Deep Learning Applications.

More problems with sampling Bayesian Estimators

Gibbs Samplers, Metropolis-Hastings

- Have a serious problem, they are difficult to parallelize making them unusable for Deep Learning Applications.

We need another way to estimate the posterior of a Bayesian method

- Variational Bayes is such alternative

More problems with sampling Bayesian Estimators

Gibbs Samplers, Metropolis-Hastings

- Have a serious problem, they are difficult to parallelize making them unusable for Deep Learning Applications.

We need another way to estimate the posterior of a Bayesian method

- Variational Bayes is such alternative

This is done by approximate $p(x|z)$ by the use of a tractable $q(x)$

- This is done by the use a trick in Variational Bayes.

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part

3 Variational Autoencoders

- Introduction
- **The Variational Bound**
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

We start with something simple

The classic

- The marginal likelihood is composed of a sum over the marginal likelihoods of individual datapoints:

$$\log p_{\theta}(\mathbf{x}_1, \dots, \mathbf{x}_N) = \sum_{i=1}^N \log p_{\theta}(\mathbf{x}_i)$$

We start with something simple

The classic

- The marginal likelihood is composed of a sum over the marginal likelihoods of individual datapoints:

$$\log p_{\theta}(\mathbf{x}_1, \dots, \mathbf{x}_N) = \sum_{i=1}^N \log p_{\theta}(\mathbf{x}_i)$$

Therefore, if we can maximize each $\log p_{\theta}(\mathbf{x}_i)$

- We maximize all $\log p_{\theta}(\mathbf{x}_1, \dots, \mathbf{x}_N)$

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part

3 Variational Autoencoders

- Introduction
- **The Variational Bound**
 - **The Kullback-Leibler Divergency**
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

The Variational Bound using the Kullback-Leibler (KL) Divergence

The best Variational Bayes approximation $q^* \in \mathcal{Q}$ is found by minimizing against the true $p(x)$

$$q^* = \arg \min_{q^* \in \mathcal{Q}} \left\{ KL(q \| p) = \int q(x) \log \left(\frac{q(x)}{p(x)} \right) dx \right\}$$

Properties of KL Divergence

Non-Negativity

- We have that $KL(q \| p) \geq 0$
- Equality holds if and only if : $q = p$ almost surely (i.e., the distributions are identical).

Properties of KL Divergence

Non-Negativity

- We have that $KL(q \| p) \geq 0$
- Equality holds if and only if : $q = p$ almost surely (i.e., the distributions are identical).

Convexity in q

- The KL divergence is convex in q (for fixed p).
- This ensures that the optimization problem in variational inference (minimizing KL divergence) has a unique global minimum, making it tractable for gradient-based methods.

Furthermore

We actually cannot minimize the KL divergence exactly,

- but we can minimize a function that is equal to it up to a constant

Furthermore

We actually cannot minimize the KL divergence exactly,

- but we can minimize a function that is equal to it up to a constant

An application of the Jensen's inequality for probability distributions

- When f is concave

$$f(E[x]) \geq E[f(x)]$$

Using it we get the maximizing the Evidence Lower Bound (ELBO)

Going back to the problem of maximizing $\log p_{\theta}(\mathbf{x}_1, \dots, \mathbf{x}_N)$

$$\begin{aligned}\log p(x) &= \log \int_z p(x, z) \\ &= \log \int_z p(x, z) \frac{q(z)}{q(z)} \\ &= \log \left[E_q \left[\frac{p(x, Z)}{q(z)} \right] \right] \\ &\geq E_q \left[\log \left(\frac{p(x, Z)}{q(z)} \right) \right] \\ &= E_q [\log p(x, Z)] - E_q [\log q(Z)]\end{aligned}$$

From minimization to maximization

Minimizing KL is equivalent to maximizing the lower bound on $\log p(x)$

$$ELBO(q) = E_q[\log p(x, Z)] - E_q[\log q(Z)]$$

- The Expectation Lower Bound (ELBO) which also appears at the Expectation Maximization

We can see this when looking at the Posterior probability

Note first that

$$p(z|x) = \frac{p(z, x)}{p(x)}$$

We can see this when looking at the Posterior probability

Note first that

$$p(z|x) = \frac{p(z, x)}{p(x)}$$

We can see how the ELBO is inserted into the KL divergence

$$\begin{aligned} KL(q(z) \| p(z|x)) &= E_q \left[\log \frac{q(Z)}{p(Z|x)} \right] \\ &= E_q [\log q(Z)] - E_q [\log p(Z|x)] \\ &= E_q [\log q(Z)] - E_q \left[\log \frac{p(Z, x)}{p(x)} \right] \\ &= - \left(\underbrace{E_q [\log p(Z, x)] - E_q [\log q(Z)]}_{ELBO} \right) + \log p(x) \\ &= -E_q \left[\log \frac{p(Z, x)}{q(Z)} \right] + \log p(x) \end{aligned}$$

Therefore

If we want to minimize the $KL(q(z) || p(z|x))$

- We need to maximize the term $E_q[\log p(x, Z)] - E_q[\log q(Z)]$

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part

3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- **Example, Mean Field Variational Inference**
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

Thus, we have

Something Notable

- The join likelihood is composed of a sum over the marginal likelihoods of individual datapoints

$$\log p_{\Theta}(z_1, z_2, z_3, \dots, z_N) = \sum_{i=1}^N \log p_{\Theta}(z_i)$$

Thus, we have

Something Notable

- The joint likelihood is composed of a sum over the marginal likelihoods of individual datapoints

$$\log p_{\Theta}(z_1, z_2, z_3, \dots, z_N) = \sum_{i=1}^N \log p_{\Theta}(z_i)$$

Typically, this family does not contain the true posterior

- because the hidden variables are dependent.
 - ▶ In the Gaussian mixture model all of the cluster assignments z_i are dependent on each other and the cluster locations $\mu_{1:K}$ given the data $x_{1:n}$.
 - ▶ These dependencies are often what makes the posterior difficult to work with.

Some Notation

$$p(z_j | z_{-j}, x)$$

- $p(z_j | z_1, \dots, z_{j-1}, z_j, \dots, z_m, x) = p(z_j | z_{-j}, x)$

Now, we optimize the ELBO for this factorized distribution

First, recall the chain rule and use it to decompose the joint

$$p(z_{1:m}, x_{1:n}) = p(x_{1:n}) \prod_{j=1}^m p(z_j | z_{1:j-1}, x_{1:n})$$

- Notice that the z variables can occur in any order in this chain.

Now, we optimize the ELBO for this factorized distribution

First, recall the chain rule and use it to decompose the joint

$$p(z_{1:m}, x_{1:n}) = p(x_{1:n}) \prod_{j=1}^m p(z_j | z_{1:j-1}, x_{1:n})$$

- Notice that the z variables can occur in any order in this chain.

Second, decompose the entropy of the variational distribution,

$$E[\log q(z_{1:m})] = \sum_{j=1}^m E_j[\log q(z_j)]$$

Now, we optimize the ELBO for this factorized distribution

First, recall the chain rule and use it to decompose the joint

$$p(z_{1:m}, x_{1:n}) = p(x_{1:n}) \prod_{j=1}^m p(z_j | z_{1:j-1}, x_{1:n})$$

- Notice that the z variables can occur in any order in this chain.

Second, decompose the entropy of the variational distribution,

$$E[\log q(z_{1:m})] = \sum_{j=1}^m E_j[\log q(z_j)]$$

We know from the previous

$$KL(q(z) || p(z|x)) = -E_q \left[\log \frac{p(Z,x)}{q(Z)} \right] + \log p(x)$$

$$KL(q(z) || p(z|x)) = \log p(x) - [E_q[\log p(Z, x) - \log q(Z)]]$$

We have then

We know that

$$p(z_j | z_{1:j-1}, x_{1:n}) = \frac{p(z_{1:j}, x_{1:n})}{p(z_{1:j-1}, x_{1:n})}$$

We have then

We know that

$$p(z_j | z_{1:j-1}, x_{1:n}) = \frac{p(z_{1:j}, x_{1:n})}{p(z_{1:j-1}, x_{1:n})}$$

Then if we multiply with $\frac{p(z_{1:j-1}, x_{1:n})}{p(z_{1:j-1}, x_{1:n})}$ the numerator can go out and seen as a constant

$$\mathcal{L} = \log p(x_{1:n}) - \sum_{j=1}^m \left(E \left[\log \frac{p(z_{1:j}, x_{1:n})}{p(z_{1:j-1}, x_{1:n})} \right] + \underbrace{E[p(z_{1:j-1}, x_{1:n})]}_{\text{const with respect to } z_j} - E_j[\log q(z_j)] \right)$$

Finally, we get

A final equation

$$\mathcal{L} \approx \log p(x_{1:n}) - \underbrace{\sum_{j=1}^m (E[\log p(z_j | z_{1:j-1}, x_{1:n})] - E_j[\log q(z_j)])}_{KL}$$

Now, Consider the ELBO as a function of $q(z_j)$

Employ the chain rule with the variable z_k as the last variable in the list

- This leads to the objective function

$$\begin{aligned}\mathcal{L} &= E[\log p(z_j | z_{-j}, x)] - E_j[\log q(z_j)] + \text{constant} \\ &\approx E[\log p(z_j | z_{-j}, x)] - E_j[\log q(z_j)]\end{aligned}$$

Now, Consider the ELBO as a function of $q(z_j)$

Employ the chain rule with the variable z_k as the last variable in the list

- This leads to the objective function

$$\begin{aligned}\mathcal{L} &= E[\log p(z_j | z_{-j}, x)] - E_j[\log q(z_j)] + \text{constant} \\ &\approx E[\log p(z_j | z_{-j}, x)] - E_j[\log q(z_j)]\end{aligned}$$

Write this objective as a function of $q(z_k)$

$$\mathcal{L}[q(z_j)] \approx \int q(z_j) E_{q_{-j}}[\log p(z_j | z_{-j}, x)] dz_j - \int q(z_j) \log q(z_j) dz_j$$

Therefore, we want

To maximize this quantity

$$\arg \max_{q_j} \left\{ \int q(z_j) E_{q_{-j}} [\log p(z_j | z_{-j}, x)] dz_j - \int q(z_j) \log q(z_j) dz_j \right\}$$

Now Optimize

Take the derivative with respect to $q(z_j)$

$$\frac{d\mathcal{L}[q(z_j)]}{dq(z_j)} = E_{q_{-j}}[\log p(z_j|z_{-j}, x)] - \log q(z_j) - \underbrace{\frac{q(z_j)}{q(z_j)}}_1 = 0$$

Now Optimize

Take the derivative with respect to $q(z_j)$

$$\frac{d\mathcal{L}[q(z_j)]}{dq(z_j)} = E_{q_{-j}}[\log p(z_j|z_{-j}, x)] - \log q(z_j) - \underbrace{\frac{q(z_j)}{q(z_j)}}_1 = 0$$

This (and Lagrange multipliers) leads to an coordinate ascent for $q(z_j)$

$$q^*(z_j) \propto \exp \{ E_{q_{-j}}[\log p(z_j|z_{-j}, x)] \}$$

Now Optimize

Take the derivative with respect to $q(z_j)$

$$\frac{d\mathcal{L}[q(z_j)]}{dq(z_j)} = E_{q_{-j}}[\log p(z_j|z_{-j}, x)] - \log q(z_j) - \underbrace{\frac{q(z_j)}{q(z_j)}}_1 = 0$$

This (and Lagrange multipliers) leads to an coordinate ascent for $q(z_j)$

$$q^*(z_j) \propto \exp \{ E_{q_{-j}} [\log p(z_j|z_{-j}, x)] \}$$

Since the denominator of the conditional does not depend on z_j

$$q^*(z_j) \propto \exp \{ E_{q_{-j}} [\log p(z_j, z_{-j}, x)] \}$$

Some Remarks

We have that

- This coordinate ascent procedure converges to a local maximum.
- The coordinate ascent update for $q(z_j)$ only depends on the other, fixed approximations $q(z_k), k \neq j$.

Some Remarks

We have that

- This coordinate ascent procedure converges to a local maximum.
- The coordinate ascent update for $q(z_j)$ only depends on the other, fixed approximations $q(z_k), k \neq j$.

In addition

- While this determines the optimal $q(z_j)$, it is still not specified.
- Depending on what form we use, the coordinate update $q(z_j)$ might not be easy to work with.

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part

3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- **A Recap of the Previous Ideas**
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

We have that

Thus, from $\log p_{\Theta}(z_1, z_2, z_3, \dots, z_N) = \sum_{i=1}^N \log p_{\Theta}(z_i)$, and knowing that KL is positive

$$\log p_{\Theta}(z_i) \geq - \int_{\mathcal{Z}} q_{\phi}(z) \log \left(\frac{p_{\Theta}(x, z_i)}{q_{\phi}(z)} \right) dx + \mathcal{T}(\Theta, \phi | z_i)$$

We have that

Thus, from $\log p_{\Theta}(z_1, z_2, z_3, \dots, z_N) = \sum_{i=1}^N \log p_{\Theta}(z_i)$, and knowing that KL is positive

$$\log p_{\Theta}(z_i) \geq - \int_{\mathcal{Z}} q_{\phi}(z) \log \left(\frac{p_{\Theta}(x, z_i)}{q_{\phi}(z)} \right) dx + \mathcal{T}(\Theta, \phi | z_i)$$

The First Term has

- $\int_{\mathcal{Z}} q_{\phi}(z) \log \left(\frac{p_{\Theta}(x, z_i)}{q_{\phi}(z)} \right) dx = D_{KL}(p_{\Theta}(x, z_i) || q_{\phi}(z))$ is the KL Divergence of the approximate of the true posterior

We have that

Thus, from $\log p_{\Theta}(z_1, z_2, z_3, \dots, z_N) = \sum_{i=1}^N \log p_{\Theta}(z_i)$, and knowing that KL is positive

$$\log p_{\Theta}(z_i) \geq - \int_{\mathcal{Z}} q_{\phi}(z) \log \left(\frac{p_{\Theta}(x, z_i)}{q_{\phi}(z)} \right) dx + \mathcal{T}(\Theta, \phi | z_i)$$

The First Term has

- $\int_{\mathcal{Z}} q_{\phi}(z) \log \left(\frac{p_{\Theta}(x, z_i)}{q_{\phi}(z)} \right) dx = D_{KL}(p_{\Theta}(x, z_i) || q_{\phi}(z))$ is the KL Divergence of the approximate of the true posterior

The Second Term

- $\mathcal{T}(\Theta, \phi | z_i)$ the variational lower bound on the marginal likelihood z_i

Now given that the $D_{KL}(p_{\Theta}(x, z_i) || q_{\phi}(z))$ is positive

We can say

$$\log p_{\Theta}(z_i) \geq \mathcal{T}(\Theta, \phi | z_i) = E_{q_{\phi}(x|z)} [-\log q_{\phi}(x|z) + \log p_{\Theta}(z)]$$

Now given that the $D_{KL}(p_{\Theta}(x, z_i) || q_{\phi}(z))$ is positive

We can say

$$\log p_{\Theta}(z_i) \geq \mathcal{T}(\Theta, \phi | z_i) = E_{q_{\phi}(x|z)} [-\log q_{\phi}(x|z) + \log p_{\Theta}(z)]$$

Actually is a KL on p_{Θ} and q_{ϕ}

$$-\log q_{\phi}(x|z) + \log p_{\Theta}(z) = \log \left(\frac{p_{\Theta}(z)}{q_{\phi}(x|z)} \right)$$

Now given that the $D_{KL}(p_{\Theta}(x, z_i) || q_{\phi}(z))$ is positive

We can say

$$\log p_{\Theta}(z_i) \geq \mathcal{T}(\Theta, \phi | z_i) = E_{q_{\phi}(x|z)} [-\log q_{\phi}(x|z) + \log p_{\Theta}(z)]$$

Actually is a KL on p_{Θ} and q_{ϕ}

$$-\log q_{\phi}(x|z) + \log p_{\Theta}(z) = \log \left(\frac{p_{\Theta}(z)}{q_{\phi}(x|z)} \right)$$

Actually, we can simplify this last function by dropping the term $-\log q_{\phi}(x|z)$ because is already in the first term

$$\mathcal{T}(\Theta, \phi | z_i) \propto E_{q_{\phi}(x|z)} [\log p_{\Theta}(z)]$$

Therefore, we can generate a loss function combining both terms $\mathcal{L}(\Theta, \phi|z)$

As the following give that they are KL divergences i.e. maximization of this minimize the divergences

$$\mathcal{L}(\Theta, \phi|z_i) = -KL(p_{\Theta}(x, z) || q_{\phi}(z)) + E_{q_{\phi}(x|z_i)} [\log p_{\Theta}(z_i)]$$

Therefore, we can generate a loss function combining both terms $\mathcal{L}(\Theta, \phi|z)$

As the following give that they are KL divergences i.e. maximization of this minimize the divergences

$$\mathcal{L}(\Theta, \phi|z_i) = -KL(p_{\Theta}(x, z) || q_{\phi}(z)) + E_{q_{\phi}(x|z_i)}[\log p_{\Theta}(z_i)]$$

Thus, we want to derive $\mathcal{L}(\Theta, \phi|z_i)$

- w.r.t. both the variational parameters ϕ and generative parameters Θ .

Problem

Problem is that MCMC is not enough

- The usual (naïve) Monte Carlo gradient estimator for this type of problem is:

$$\begin{aligned}\nabla_{\phi} E_{q_{\phi}(x)} (\log p_{\Theta}(z_i)) &= E \left(\nabla_{q_{\phi}(x)} q_{\phi}(x) \log p_{\Theta}(z_i) \right) \\ &\approx \frac{1}{L} \sum_{l=1}^L \nabla_{q_{\phi}(x)} q_{\phi}(x_l) \log p_{\Theta}(z_i | x_l) \\ &\approx E_{q_{\phi}(x_l)} [\log p_{\Theta}(z | x_l)]\end{aligned}$$

Problem with this first attempt of using this for maximization

Here, we have that $x \sim q_\phi(x|z_i)$, yes z_i is a hidden variable

- This gradient estimator exhibits very high variance and non practical

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part

3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- **Re-Parameterization Trick**
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

First, explicit re-parameterization Gradients

Suppose we would like to optimize an expectation $E_{q_{\phi}(x)} [f(x)]$

- A continuously differentiable function $f(z)$ w.r.t. the parameters of the distribution.

First, explicit re-parameterization Gradients

Suppose we would like to optimize an expectation $E_{q_\phi(x)} [f(x)]$

- A continuously differentiable function $f(z)$ w.r.t. the parameters of the distribution.

We assume that we can find a standardization function

$$S_\phi(x) = \epsilon \sim q(x)$$

- Such that you can do the following $x = S_\phi^{-1}(\epsilon)$

First, explicit re-parameterization Gradients

Suppose we would like to optimize an expectation $E_{q_\phi(x)} [f(x)]$

- A continuously differentiable function $f(z)$ w.r.t. the parameters of the distribution.

We assume that we can find a standardization function

$$S_\phi(x) = \epsilon \sim q(x)$$

- Such that you can do the following $x = S_\phi^{-1}(\epsilon)$

For example, for a Gaussian distribution $N(\mu, \sigma^2)$

- We can use the following function $S_{\mu,\sigma}(x) = \frac{x-\mu}{\sigma} \sim N(0, 1)$ then we have we can sample $\epsilon \sim N(0, 1)$

$$S_{\mu,\sigma}^{-1}(\epsilon) = \sigma \cdot \epsilon + \mu$$

Therefore, we can express the $E_{q_\phi(x)} [f(x)]$

As follow

$$E_{q_\phi(x)} [f(x)] = E_{q(\epsilon)} [f(S_\phi^{-1}(\epsilon))]$$

Therefore, we can express the $E_{q_\phi(x)} [f(x)]$

As follow

$$E_{q_\phi(x)} [f(x)] = E_{q(\epsilon)} \left[f \left(S_\phi^{-1}(\epsilon) \right) \right]$$

This allows us to compute the gradient of the expectation as the expectation of the gradients by rule chain:

$$\begin{aligned} \nabla_\phi E_{q_\phi(x)} [f(x)] &= E_{q(\epsilon)} \left[\nabla_\phi f \left(S_{\mu,\sigma}^{-1}(\epsilon) \right) \right] \\ &= E_{q(\epsilon)} \left[\nabla_z f \left(S_{\mu,\sigma}^{-1}(\epsilon) \right) \nabla_\phi S_\phi^{-1}(\epsilon) \right] \end{aligned}$$

Therefore

An alternative was proposed to avoid the inverse of $S_\phi(x)$

- For this, we first we express the classic stuff

$$\begin{aligned}\nabla_\phi E_{q_\phi(x)}[f(x)] &= E_{q_\phi(x)}[\nabla_x f(x) \nabla_\phi x] \\ \nabla_\phi x &= \nabla_\phi S_\phi^{-1}(\epsilon) \Big|_{\epsilon=S_\phi(x)}\end{aligned}$$

Therefore

An alternative was proposed to avoid the inverse of $S_\phi(x)$

- For this, we first we express the classic stuff

$$\begin{aligned}\nabla_\phi E_{q_\phi(x)}[f(x)] &= E_{q_\phi(x)}[\nabla_x f(x) \nabla_\phi x] \\ \nabla_\phi x &= \nabla_\phi S_\phi^{-1}(\epsilon) \Big|_{\epsilon=S_\phi(x)}\end{aligned}$$

It was proposed to compute $\nabla_\phi x$ by implicit differentiation

- Implicit differentiation makes use of the chain rule to differentiate a function which cannot be explicitly expressed in the form $y = f(x)$

$$\frac{df(y(x))}{dx} = \frac{df}{dy} \times \frac{dy}{dx}$$

Then

We apply the total gradient to the equality $S_\phi(x) = \epsilon$

- We finish with the following given that $S_\phi(x)$ depends on ϕ directly by the subscript and from x indirectly. Thus, $x = x(\phi)$:

$$\frac{dS_\phi(x)}{dx} \times \frac{dx}{d\phi} + \frac{dS_\phi(x)}{d\phi} = 0$$

Then, we have that

Therefore, we have that

$$\frac{dx}{d\phi} = - \left(\frac{dS_{\phi}(x)}{dx} \right)^{-1} \times \frac{dS_{\phi}(x)}{d\phi}$$

Then, we have that

Therefore, we have that

$$\frac{dx}{d\phi} = - \left(\frac{dS_{\phi}(x)}{dx} \right)^{-1} \times \frac{dS_{\phi}(x)}{d\phi}$$

This expression for the gradient only requires differentiating the standardization function

- Basically we get the value and invert it, we do not need an explicit version of the inverse function, then derive it.

Example

Univariate Normal distribution $N(\mu, \sigma^2)$

- We have the standardization function $S_{\mu, \sigma}(x) = \frac{x - \mu}{\sigma} = \epsilon \sim N(0, 1)$

Example

Univariate Normal distribution $N(\mu, \sigma^2)$

- We have the standardization function $S_{\mu, \sigma}(x) = \frac{x - \mu}{\sigma} = \epsilon \sim N(0, 1)$

We have then, if we use the derivatives

$$\frac{dx}{d\mu} = - \left(\frac{dS_{\mu, \sigma}(x)}{dz} \right)^{-1} \times \frac{dS_{\mu, \sigma}(x)}{d\mu} = - \frac{-\frac{1}{\sigma}}{\frac{1}{\sigma}} = 1$$

$$\frac{dx}{d\sigma} = - \left(\frac{dS_{\mu, \sigma}(x)}{dz} \right)^{-1} \times \frac{dS_{\mu, \sigma}(x)}{d\sigma} = \frac{-\frac{(z - \mu)}{\sigma^2}}{\frac{1}{\sigma}} = \frac{(z - \mu)}{\sigma} \sim N(0, 1)$$

If we compare against inverting and deriving

We have the following inverse of the function

$$z = S_{\mu, \sigma}^{-1}(x) = \mu + \sigma \epsilon$$

If we compare against inverting and deriving

We have the following inverse of the function

$$z = S_{\mu, \sigma}^{-1}(x) = \mu + \sigma \epsilon$$

Then, we have that with $\epsilon \sim N(0, 1)$

$$\frac{dz}{d\mu} = 1,$$

$$\frac{dz}{d\sigma} = \epsilon$$

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part

3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- **Re-Parameterization Trick**
 - **Now the application for the Variational Problem**
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

We have that

The essential parameterization trick is quite simple

- Let x be a continuous random variable and $x \sim q_\phi(x|z)$ a conditional distribution

We have that

The essential parameterization trick is quite simple

- Let x be a continuous random variable and $x \sim q_\phi(x|z)$ a conditional distribution

It is possible to extend $S_\phi(x) = \epsilon$

$$g_{\phi,z}(x) = \epsilon$$

- with $\epsilon \sim p(x)$ and $g_\phi(\cdot)$ is some vector valued function parameterized by ϕ, z

This parameterization is useful for our case

It can be used to rewrite an expectation w.r.t $q_\phi(x|z)$

- Such that the Monte Carlo estimate of the expectation is differentiable w.r.t. ϕ

This parameterization is useful for our case

It can be used to rewrite an expectation w.r.t $q_\phi(x|z)$

- Such that the Monte Carlo estimate of the expectation is differentiable w.r.t. ϕ

We want the following

- Given the mapping $g_{\phi,z}(x) = \epsilon$ (Here, x is a vector and $dx = \prod_i dx_i$ are infinitesimals)

$$q_\phi(x|z) \approx g_{\phi,z}(x)$$

Now, we can do the following

Given that the error is a small variation, we can say that in an small area

$$q_{\phi}(x|z) \prod_i dx_i = p(x) \prod_i d\epsilon_i$$

Now, we can do the following

Given that the error is a small variation, we can say that in an small area

$$q_{\phi}(x|z) \prod_i dx_i = p(x) \prod_i d\epsilon_i$$

Therefore

$$\int q_{\phi}(x|z) f(x) dx = \int p(x) f(x) dx = \int p(x) f(g_{\phi,z}(x)) dx$$

We can then construct then a variational estimation

We have that

$$\int q_{\phi}(x|z) f(x) dx = \frac{1}{L} \sum_{l=1}^L f(g_{\phi,z}(x)) \text{ with } \epsilon_l \sim p(x)$$

Therefore, we have

Using the previous ideas

$$\nabla_{\phi} E_{q_{\phi}(x)} [f(g_{\phi,z}(x))] = E_{q_{\phi}(x)} \left[\nabla_x f(g_{\phi,z}(x)) \frac{dx}{d\phi, z} \right]$$

$$\frac{dx}{d\phi, z} = - \left(\frac{dg_{\phi,z}(x)}{dx} \right)^{-1} \times \frac{dg_{\phi,z}(x)}{d\phi, z}$$

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part

3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- **Stochastic Gradient Variational Bayes (SGVB)**
- Example, Variational Autoencoders as Generative Models
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

Therefore

We have that

$$\tilde{\mathcal{L}}^A(\Theta, \phi|z_i) \approx \mathcal{L}(\Theta, \phi|z_i)$$

Therefore

We have that

$$\tilde{\mathcal{L}}^A(\Theta, \phi|z_i) \approx \mathcal{L}(\Theta, \phi|z_i)$$

Thus, we use the estimation with $x_{i,l} = g_\phi(\epsilon_{i,l})$ and $\epsilon_l \sim p(\epsilon)$

$$\tilde{\mathcal{L}}^A(\Theta, \phi|z_i) = \frac{1}{L} \sum_{l=1}^L [\log p_\Theta(z_i, x_{i,l}) - \log q_\phi(x_{i,l}|z_i)]$$

Therefore

We have that

$$\tilde{\mathcal{L}}^A(\Theta, \phi|z_i) \approx \mathcal{L}(\Theta, \phi|z_i)$$

Thus, we use the estimation with $x_{i,l} = g_\phi(\epsilon_{i,l})$ and $\epsilon_l \sim p(\epsilon)$

$$\tilde{\mathcal{L}}^A(\Theta, \phi|z_i) = \frac{1}{L} \sum_{l=1}^L [\log p_\Theta(z_i, x_{i,l}) - \log q_\phi(x_{i,l}|z_i)]$$

Actually, if $p(x)$ is seen as a normalization factor

$$\log \left(\frac{p_\Theta(z|x)}{q_\phi(x)} \right) = \log \left(\frac{p_\Theta(z, x)}{p(z) q_\phi(x)} \right) \approx \log \left(\frac{p_\Theta(z, x)}{q_\phi(x|z)} \right)$$

In this was, we have

We can say that

$$\log \left(\frac{p_{\Theta}(z_i, x_{i,l})}{q_{\phi}(x_{i,l}|z_i)} \right) \approx \log p_{\Theta}(z_i, x_{i,l}) - \log q_{\phi}(x_{i,l}|z_i)$$

Final Algorithm

Minibatch version of the Auto-Encoding VB (AEVB) algorithm

- Initialize Parameters Θ, ϕ
- Repeat
 - $X^M \leftarrow$ Random Mini-batch of M data points
 - $\epsilon \leftarrow$ Random samples from noise distribution $p(\epsilon)$
 - $g \leftarrow \nabla_{\Theta, \phi} \tilde{\mathcal{L}}^A(\Theta, \phi | z_i)$ (Gradient Mini-batch estimator)
 - Θ, ϕ update parameters using the gradient
- return Θ, ϕ

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part

3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- **Example, Variational Autoencoders as Generative Models**
- Autoencoder Applications
 - Generative Models
 - CNN Variational Autoencoder

The original variational auto-encoder

It is a continuous latent variable model.

- The model is intended to learn a latent space $\mathcal{Z} = \mathbb{R}^t$ using a given set of samples $\{x_n\} \subseteq \mathcal{Y} = \mathbb{R}^d$
- Such that $t \ll d$

The original variational auto-encoder

It is a continuous latent variable model.

- The model is intended to learn a latent space $\mathcal{Z} = \mathbb{R}^t$ using a given set of samples $\{x_n\} \subseteq \mathcal{Y} = \mathbb{R}^d$
- Such that $t \ll d$

Therefore

- The model consists of the generative model $p(x|z)$ given a fixed prior $p(z)$

Bernoulli MLP as decoder

In this case, let $p_{\Theta}(x|z)$ be a multivariate Bernoulli

- The probability is calculated using a fully connected neural network with a single layer

$$\log p(x|z) = \sum_{i=1}^D x_i \log y_i + (1 - x_i) \log (1 - y_i)$$

- ▶ where $\mathbf{y} = f_{\sigma}(W_2 \tanh(W_1 z + b_1) + b_2)$ with f_{σ} is the sigmoidal function
- ▶ with $\Theta = \{W_1, W_2, b_1, b_2\}$ are the weight and biases of the MLP

Gaussian MLP as encoder

We have the following function

$$\log p(z|x) = \log N(\mu, \sigma^2 I)$$

$$\text{where } \mu = W_4 h + b_4$$

$$\log \sigma^2 = W_5 h + b_5$$

$$h = \tanh(W_3 x + b_3)$$

We have then

Let the prior over isotropic multivariate Gaussian

$$p_{\Theta}(z) = N(0, I)$$

- Note that in this case, the prior lacks parameters.

We have then

Let the prior over isotropic multivariate Gaussian

$$p_{\Theta}(z) = N(0, I)$$

- Note that in this case, the prior lacks parameters.

We let $p_{\Theta}(x|z)$

- Bernoulli whose distribution parameters are computed from z with a MLP

We have then

Let the prior over isotropic multivariate Gaussian

$$p_{\Theta}(z) = N(0, I)$$

- Note that in this case, the prior lacks parameters.

We let $p_{\Theta}(x|z)$

- Bernoulli whose distribution parameters are computed from z with a MLP

Note

- Note the true posterior $p_{\Theta}(z|x)$ is in this case intractable.

Now for our estimation

We have then the following

$$\log q_{\phi}(z|x) = \log N(\mu, \sigma^2 I)$$

- where the mean and s.d. of the approximate posterior, μ and σ , are outputs of the encoding MLP, i.e. nonlinear functions of data point x and the variational parameters ϕ

Now for our estimation

We have then the following

$$\log q_{\phi}(z|x) = \log N(\mu, \sigma^2 I)$$

- where the mean and s.d. of the approximate posterior, μ and σ , are outputs of the encoding MLP, i.e. nonlinear functions of data point x and the variational parameters ϕ

Then, we sample $z_{i,l} \sim q_{\phi}(z|x_i)$ using the fact that we have a diagonal covariance $\sigma^2 I$

$$z_{i,l} = \mu_i + \sigma_i \odot \epsilon_l \text{ and } \epsilon_l \sim N(0, I)$$

Then, we have that

In this model both

- $p_{\Theta}(z)$ and $q_{\phi}(z|x_i)$ are Gaussian

Then, we have that

In this model both

- $p_{\Theta}(z)$ and $q_{\phi}(z|x_i)$ are Gaussian

The resulting estimator for this model comes from the classic where the l is the meaning of multiple samples of z_i

$$\mathcal{L}(\Theta, \phi|x_i) = D_{KL}(q_{\phi}(z|x_i) || p_{\Theta}(z)) + \underbrace{\frac{1}{L} \sum_{l=1}^L -\log p_{\theta}(x_i|z_{i,l})}_{E[-\log p_{\theta}(x_i|z_i)]}$$

- Basically the negative of $-\int_{\mathcal{Z}} q_{\phi}(z) \log \left(\frac{p_{\Theta}(x, z_i)}{q_{\phi}(z)} \right) dz + \mathcal{T}(\Theta, \phi|z_i)$ which we want to maximize but when minimizing we take the negative.

Therefore

We have that $q_{\phi}(z|x_i) = N(z|\mu(x), \sigma^2(x) I)$

- Where at the encoder

$$\mu = W_4 x + b_4$$

$$\log \sigma^2 = W_5 x + b_5$$

Therefore

We have that $q_\phi(z|x_i) = N(z|\mu(x), \sigma^2(x)I)$

- Where at the encoder

$$\begin{aligned}\mu &= W_4 x + b_4 \\ \log \sigma^2 &= W_5 x + b_5\end{aligned}$$

The prior of $p(z) \approx N(z|0, I)$

- Then, the KL divergence is

$$D_{KL}(q_\phi(z|x) || p_\Theta(z)) = E_{q_\phi(z|x)} \left[\log \frac{q_\phi(z|x)}{p_\Theta(z)} \right]$$

Therefore

We have that $q_\phi(z|x_i) = N(z|\mu(x), \sigma^2(x)I)$

- Where at the encoder

$$\begin{aligned}\mu &= W_4 x + b_4 \\ \log \sigma^2 &= W_5 x + b_5\end{aligned}$$

The prior of $p(z) \approx N(z|0, I)$

- Then, the KL divergence is

$$D_{KL}(q_\phi(z|x) || p_\Theta(z)) = E_{q_\phi(z|x)} \left[\log \frac{q_\phi(z|x)}{p_\Theta(z)} \right]$$

Expand the logarithm

$$\log \frac{q_\phi(z|x)}{p_\Theta(z)} = \log q_\phi(z|x) - \log p_\Theta(z)$$

Given that we have Gaussian's

Assume the encoder outputs a diagonal Gaussian distribution

$$\log q_{\phi}(z|x) = -\frac{1}{2} \left(d \log 2\pi + d \log \sigma^2 + \frac{\|z - \mu\|^2}{\sigma^2} \right)$$
$$p_{\Theta}(z) = -\frac{1}{2} z^T z - \frac{d}{2} \log 2\pi$$

Given that we have Gaussian's

Assume the encoder outputs a diagonal Gaussian distribution

$$\log q_{\phi}(z|x) = -\frac{1}{2} \left(d \log 2\pi + d \log \sigma^2 + \frac{\|z - \mu\|^2}{\sigma^2} \right)$$
$$p_{\Theta}(z) = -\frac{1}{2} z^T z - \frac{d}{2} \log 2\pi$$

Putting all together, we have

$$D_{KL} \left(q_{\phi}(z|x_i) || p_{\Theta}(z) \right) = E_{q_{\phi}(z|x_i)} \left[-\frac{1}{2} d \log 2\pi - \frac{1}{2} d \log \sigma^2 - \frac{1}{2} \frac{\|z - \mu\|^2}{\sigma^2} + \frac{1}{2} z^T z + \frac{d}{2} \log 2\pi \right]$$

We have the following equality's on the Expectation

We have the following

$$\begin{aligned} E_{q_\phi(z|x_i)} [\|z - \mu\|^2] &= E_{q_\phi(z|x_i)} [\text{tr}([z - \mu][z - \mu]^T)] = \text{tr}(\sigma^2 I) = d\sigma^2 \\ E_{q_\phi(z|x_i)} [z^T z] &= \mu^T \mu + \text{tr}(\sigma I) = \mu^T \mu + d\sigma^2 \end{aligned}$$

Therefore, we apply the expected value

We have the following

$$\begin{aligned} D_{KL} \left(q_{\phi} (z|x_i) || p_{\Theta} (z) \right) &= -\frac{1}{2} \log 2\pi - \frac{1}{2} \log \sigma^2 - \frac{1}{2} \frac{d\sigma^2}{\sigma^2} + \frac{1}{2} \mu^T \mu + \frac{1}{2} d\sigma^2 + \frac{d}{2} \log 2\pi \\ &\approx -\frac{1}{2} d \log \sigma^2 - \frac{d}{2} + \frac{1}{2} \sum_{i=1}^d \mu_i^2 + \frac{1}{2} d\sigma^2 \end{aligned}$$

Therefore, we apply the expected value

We have the following

$$\begin{aligned} D_{KL} (q_\phi (z|x_i) || p_\Theta (z)) &= -\frac{1}{2} \log 2\pi - \frac{1}{2} \log \sigma^2 - \frac{1}{2} \frac{d\sigma^2}{\sigma^2} + \frac{1}{2} \mu^T \mu + \frac{1}{2} d\sigma^2 + \frac{d}{2} \log 2\pi \\ &\approx -\frac{1}{2} d \log \sigma^2 - \frac{d}{2} + \frac{1}{2} \sum_{i=1}^d \mu_i^2 + \frac{1}{2} d\sigma^2 \end{aligned}$$

Which can be seen as with $\sigma_i^2 = \sigma^2$

$$D_{KL} (q_\phi (z|x_i) || p_\Theta (z)) = -\frac{1}{2} \sum_{i=1}^d \left(\log \sigma_i^2 + 1 - \frac{1}{2} \mu_i^2 + \sigma_i^2 \right)$$

Therefore, we apply the expected value

We have the following

$$\begin{aligned} D_{KL} (q_{\phi} (z|x_i) || p_{\Theta} (z)) &= -\frac{1}{2} \log 2\pi - \frac{1}{2} \log \sigma^2 - \frac{1}{2} \frac{d\sigma^2}{\sigma^2} + \frac{1}{2} \mu^T \mu + \frac{1}{2} d\sigma^2 + \frac{d}{2} \log 2\pi \\ &\approx -\frac{1}{2} d \log \sigma^2 - \frac{d}{2} + \frac{1}{2} \sum_{i=1}^d \mu_i^2 + \frac{1}{2} d\sigma^2 \end{aligned}$$

Which can be seen as with $\sigma_i^2 = \sigma^2$

$$D_{KL} (q_{\phi} (z|x_i) || p_{\Theta} (z)) = -\frac{1}{2} \sum_{i=1}^d \left(\log \sigma_i^2 + 1 - \frac{1}{2} \mu_i^2 + \sigma_i^2 \right)$$

Finally, what we want to minimize this the ELBO will be maximized

$$\mathcal{L} (\Theta, \phi | x_i) = \underbrace{-\frac{1}{2} \sum_{i=1}^d \left(1 + \log (\sigma_i^2) - (\mu_i)^2 - (\sigma_i)^2 \right)}_{D_{KL} (q_{\phi} (z|x_i) || p_{\Theta} (z))} + \underbrace{\frac{1}{L} \sum_{l=1}^L -\log p_{\theta} (x_i | z_{i,l})}_{E[-\log p_{\theta} (x|z)]}$$

Now what about the real Loss function?

We can do the following if we think that $E [\log p_{\theta} (x|z)]$ is the reconstruction log of the data i.e $x \sim p_{\theta} (x|z)$

- We can see that $p_{\theta} (x|z)$, if it is modeled as a Bernoulli distribution because you rebuild it or not

$$p_{\theta} (x|z) = p^x (1 - p)^{1-x}$$

Now what about the real Loss function?

We can do the following if we think that $E[\log p_\theta(x|z)]$ is the reconstruction log of the data i.e $x \sim p_\theta(x|z)$

- We can see that $p_\theta(x|z)$, if it is modeled as a Bernoulli distribution because you rebuild it or not

$$p_\theta(x|z) = p^x (1 - p)^{1-x}$$

Where we use $p = \sigma(z)$ the output of the decoder

$$p_\theta(x|z) = \sigma(z)^x (1 - \sigma(z))^{1-x} \Rightarrow -\log p_\theta(x|z) = -x \log(\sigma(z)) - (1 - x) \log(1 - \sigma(z))$$

Now what about the real Loss function?

We can do the following if we think that $E[\log p_\theta(x|z)]$ is the reconstruction log of the data i.e $x \sim p_\theta(x|z)$

- We can see that $p_\theta(x|z)$, if it is modeled as a Bernoulli distribution because you rebuild it or not

$$p_\theta(x|z) = p^x (1 - p)^{1-x}$$

Where we use $p = \sigma(z)$ the output of the decoder

$$p_\theta(x|z) = \sigma(z)^x (1 - \sigma(z))^{1-x} \Rightarrow -\log p_\theta(x|z) = -x \log(\sigma(z)) - (1 - x) \log(1 - \sigma(z))$$

Therefore when you take a minibatch or the $E[\log p_\theta(x_i|z_i)]$

$$E[-\log p_\theta(x|z)] = BCE[\sigma(z), x]$$

Therefore the final Loss function is

Two parts the divergence and the binary cross entropy

$$\mathcal{L}(\Theta, \phi | x_i) = \underbrace{-\frac{1}{2} \sum_{i=1}^d \left(1 + \log(\sigma_i^2) - (\mu_i)^2 - (\sigma_i)^2 \right)}_{D_{KL}(q_\phi(z|x_i) || p_\Theta(z))} + \underbrace{BCE[\sigma(z), x]}_{E[-\log p_\theta(x|z)]}$$

Therefore the final Loss function is

Two parts the divergence and the binary cross entropy

$$\mathcal{L}(\Theta, \phi | x_i) = \underbrace{-\frac{1}{2} \sum_{i=1}^d \left(1 + \log(\sigma_i^2) - (\mu_i)^2 - (\sigma_i)^2 \right)}_{D_{KL}(q_\phi(z|x_i) || p_\Theta(z))} + \underbrace{BCE[\sigma(z), x]}_{E[-\log p_\theta(x|z)]}$$

A last comment is the following, this model is not exactly a generative model but

- the $q(z|x)$ (simple and tractable posteriors) must close enough to $N(0, I)$.
- We can sample from $N(0, I)$ to get a input for the decoder and get some generated sample.

Given the inception of the Variational part

It is coming from the work of [7]

- Bishop, C.M.: Pattern Recognition and Machine Learning (Information Science and Statistics). Springer-Verlag, Berlin, Heidelberg (2006)

Given the inception of the Variational part

It is coming from the work of [7]

- Bishop, C.M.: Pattern Recognition and Machine Learning (Information Science and Statistics). Springer-Verlag, Berlin, Heidelberg (2006)

Variational Autoencoder (VAE)

- VAE are generative models that attempt to describe data generation through a probabilistic distribution.

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part

3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- **Autoencoder Applications**
 - Generative Models
 - CNN Variational Autoencoder

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part

3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- **Autoencoder Applications**
 - **Generative Models**
 - CNN Variational Autoencoder

Generative Model

At the encoder part

- The posterior distribution $q_{\phi}(x|z_i)$ which is derived by the encoder

Generative Model

At the encoder part

- The posterior distribution $q_{\phi}(x|z_i)$ which is derived by the encoder

Thus, we have

- This is regularized towards a continuous and complete distribution in the shape of the predefined prior of the latent variables $p_{\Theta}(x)$.

Therefore

Once trained

- One can simply samples random variables from the the same prior, and feed it to the decoder.

Therefore

Once trained


- One can simply sample random variables from the the same prior, and feed it to the decoder.

Since the decoder was trained generate x from $p_{\Theta}(x_i|z)$

- It would generate a meaningful generated sample

Example

Original Images vs the generated



Sample from the original MNIST dataset



VAE generated MNIST IMAGES

Outline

1 The Idea of The Autoencoder

- Introduction
- The Bottleneck Idea
- Training Autoencoders
- Encoder/Decoder Capacity
- Right Autoencoder Design: Use regularization
- Autoencoders as an initialization method

2 Types of Autoencoders

- Sparse Autoencoders
- Denoising Autoencoders
- Contractive Autoencoders
- Example, Architecture of the U-Net
 - Encoder Part

3 Variational Autoencoders

- Introduction
- The Variational Bound
 - The Kullback-Leibler Divergency
- Example, Mean Field Variational Inference
- A Recap of the Previous Ideas
- Re-Parameterization Trick
 - Now the application for the Variational Problem
- Stochastic Gradient Variational Bayes (SGVB)
- Example, Variational Autoencoders as Generative Models
- **Autoencoder Applications**
 - Generative Models
 - **CNN Variational Autoencoder**

Consider N images $\{X_l\}_{l=1}^N$, with $X_l \in \mathbb{R}^{N_x \times N_y \times C}$

We introduce the decoder

- To introduce the image decoder (generative model) in its simplest form, we first consider a decoder with $L = 2$ layers.

Code Entry \Downarrow

$$\text{Layer 2: } \tilde{S}^{n,2} = \sum_{k_2=1}^{K_2} D^{k_2,2} * S^{n,k_2,2} \quad (1)$$

$$\text{Unpool : } S^{n,1} = \text{unpool}(\tilde{S}^{n,2}) \quad (2)$$

$$\text{Layer 1: } \tilde{S}^{n,1} = \sum_{k_1=1}^{K_1} D^{k_1,2} * S^{n,k_1,2} \quad (3)$$

$$\text{Data Generation: } X_l \sim N(\tilde{S}^{n,1}, \alpha_0^{-1} I) \quad (4)$$

Notation

The following are 3D tensors

- $D^{k_l, l}$
- $S^{n, l}$
- $\tilde{S}^{n, l}$

Notation

The following are 3D tensors

- $D^{k_l, l}$
- $S^{n, l}$
- $\tilde{S}^{n, l}$

2D Activation Maps

- $S^{n, k, l}$ as slices of 3D tensors $S^{n, l}$

Notation

The following are 3D tensors

- $D^{k_l, l}$
- $S^{n, l}$
- $\tilde{S}^{n, l}$

2D Activation Maps

- $S^{n, k, l}$ as slices of 3D tensors $S^{n, l}$

Finally

- $D^{k, l} * S^{n, k, l} =$ each of the K_{l-1} “slices” of $D^{k_l, l}$ is convolved with the spatially-dependent $S^{n, k, l}$

Meaning for equation (4)

It indicates

- 1 $E[X_l] = \tilde{S}^{n,1}$
- 2 $X_l - E(X_l)$ is iid zero mean Gaussina with preccision α_0

The Stochastic Unpooling

$S^{n,k,l}$ is partitioned into contiguous $p_x \times p_y$ pooling blocks

- $z_{i,j}^{n,k,1} \in \{0, 1\}^{p_x p_y}$ be a vector of $p_x p_y - 1$ zeros and a single one.

The Stochastic Unpooling

$S^{n,k_l,l}$ is partitioned into contiguous $p_x \times p_y$ pooling blocks

- $z_{i,j}^{n,k,1} \in \{0, 1\}^{p_x p_y}$ be a vector of $p_x p_y - 1$ zeros and a single one.

$z_{i,j}^{n,k,1}$ corresponds to pooling block (i, j) in $S^{n,k_1,1}$

- The location of the non-zero element of $z_{i,j}^{n,k,1}$ identifies the location of the single non-zero element in the corresponding pooling block of $S^{n,k_1,1}$.

Therefore

The non-zero element in pooling block (i, j)

- $S^{n,k_1,1}$ is set to $\tilde{S}_{i,j}^{n,k_1,2}$

Therefore

The non-zero element in pooling block (i, j)

- $S^{n,k_1,1}$ is set to $\tilde{S}_{i,j}^{n,k_1,2}$

Thus, we have the following multinomial for the unpooling

$$z_{i,j}^{n,k_1,1} \sim \text{Mult} \left(1; \frac{1}{p_x p_y}, \dots, \frac{1}{p_x p_y} \right)$$

Finally, the Encoder

Something Notable

Code Entry \Downarrow

$$\text{Layer 1: } \tilde{C}^{n,k_1,1} = X_l *_s F^{k_1,1} \quad (5)$$

$$\text{Pool : } C^{n,1} \sim \text{pool} \quad (6)$$

$$\text{Layer 2: } \tilde{C}^{n,k_2,2} = C^{n,1} *_s F^{k_2,2} \quad (7)$$

$$\text{Data Generation: } s_n \sim N \left(\mu_0 \tilde{C}^{n,2}, \text{diag} \left(\sigma_\phi \left[\tilde{C}^{n,2} \right] \right) \right) \quad (8)$$

This was trained in a semisupervised way

By using a loss function

$$\begin{aligned}\mathcal{L}_{\phi,\alpha,\psi}(X,Y) = & \xi \left\{ E_{q_{\phi}(s|X)} [\log p_{\psi}(Y|s)] \right\} \\ & + \underbrace{E_{q_{\phi}(s,z|X)} [\log p_{\alpha}(X,s,z) - \log q_{\phi}(s,z|X)]}_{\mathcal{U}_{\phi,\alpha}(X)}\end{aligned}$$

This was trained in a semisupervised way

By using a loss function

$$\mathcal{L}_{\phi,\alpha,\psi}(X, Y) = \xi \left\{ E_{q_{\phi}(s|X)} [\log p_{\psi}(Y|s)] \right\} \\ + \underbrace{E_{q_{\phi}(s,z|X)} [\log p_{\alpha}(X, s, z) - \log q_{\phi}(s, z|X)]}_{\mathcal{U}_{\phi,\alpha}(X)}$$

The lower bound for the entire dataset is then

$$\mathcal{J}_{\phi,\alpha,\psi} = \sum_{(X,Y) \in \mathcal{D}_c} \mathcal{L}_{\phi,\alpha,\psi}(X, Y) + \sum_{X \in \mathcal{D}_u} \mathcal{U}_{\phi,\alpha}(X)$$

Benchmark for Classic

We have

Table 1: Classification error (%) and testing time (ms per image) on benchmarks.

Method	MNIST		CIFAR-10		CIFAR-100		Caltech 101		Caltech 256	
	test error	test time	test error	test time	test error	test time	test error	test time	test error	test time
Gibbs [8]	0.37	3.1	8.21	10.4	34.33	10.4	12.87	50.4	29.50	52.3
MCEM [8]	0.45	0.8	9.04	1.1	35.92	1.1	13.51	8.8	30.13	8.9
VAE-d	0.42	0.007	10.74	0.02	37.96	0.02	14.79	0.3	32.18	0.3
VAE (Ours)	0.38	0.007	8.19	0.02	35.01	0.02	11.99	0.3	29.33	0.3

Method	ImageNet 2012			ImageNet Pretrained for			
	top-1 error	top-5 error	test time	Caltech 101		Caltech 256	
				test error	test time	test error	test time
MCEM [8]	37.9	16.1	14.4	6.85	14.1	22.10	14.2
VAE (Ours)	38.2	15.7	1.0	6.91	0.9	22.53	0.9

Other examples

We have a long list

- Autoencoders for classification
- Autoencoders for clustering
- Autoencoders for anomaly detection
- Autoencoders for recommendation systems



D. Bank, N. Koenigstein, and R. Giryes, “Autoencoders,” *CoRR*, vol. abs/2003.05991, 2020.



P. Baldi, “Autoencoders, unsupervised learning, and deep architectures,” in *Proceedings of ICML Workshop on Unsupervised and Transfer Learning* (I. Guyon, G. Dror, V. Lemaire, G. Taylor, and D. Silver, eds.), vol. 27 of *Proceedings of Machine Learning Research*, (Bellevue, Washington, USA), pp. 37–49, PMLR, 02 Jul 2012.



P. Baldi and K. Hornik, “Neural networks and principal component analysis: Learning from examples without local minima,” *Neural networks*, vol. 2, no. 1, pp. 53–58, 1989.



S. J. Pan and Q. Yang, “A survey on transfer learning,” *IEEE Transactions on knowledge and data engineering*, vol. 22, no. 10, pp. 1345–1359, 2009.



O. Ronneberger, P. Fischer, and T. Brox, “U-net: Convolutional networks for biomedical image segmentation,” in *International Conference on Medical image computing and computer-assisted intervention*, pp. 234–241, Springer, 2015.



D. P. Kingma and M. Welling, “Auto-encoding variational bayes,” 2013.



C. M. Bishop, *Pattern Recognition and Machine Learning (Information Science and Statistics)*.
Secaucus, NJ, USA: Springer-Verlag New York, Inc., 2006.